

# hw4\_p3a

June 3, 2019

```
[9]: %matplotlib inline

[10]: import torch
import torchvision
import torchvision.transforms as transforms

[11]: # LOADS CIFAR10 images which are 32 x 32 x 3 RGB images
# iter(trainloader/testloader) are iterables that come in pairs (image, label)

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

```
[12]: import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
```

```

        # in features --> flattened 32 x 32 x 3 image, out features = label
        → (0,1,..., 9), use bias
        self.linear = nn.Linear(3072, 10, bias=True)

    def forward(self, x):
        x = x.view(4, -1) # flatten the image 4 x 32 x 32 x 3 --> 4 x 3072 x 1
        return self.linear(x)

net = Net()

```

[ ]:

## 1 Choose Hyperparameters

```

[13]: num_epochs = 12 # number of epochs to train for
      momentum = 0.6 # momentum for Stochastic Gradient Descent
      lr = 0.0001 # learning rate (eta) for gradient descent

```

## 2 3. Define a Loss function and optimizer

Let's use a Classification Cross-Entropy loss and SGD with momentum.

```

[14]: import torch.optim as optim

      criterion = nn.CrossEntropyLoss()
      optimizer = optim.SGD(net.parameters(), lr=lr, momentum=momentum)

```

## 3 4. Train the network

This is when things start to get interesting. We simply have to loop over our data iterator, and feed the inputs to the network and optimize.

```

[15]: def calc_accuracy(dataloader):
      correct = 0
      total = 0
      with torch.no_grad():
          for data in dataloader:
              images, labels = data
              outputs = net(images)
              _, predicted = torch.max(outputs.data, 1)
              total += labels.size(0)
              correct += (predicted == labels).sum().item()

      return 100.0 * correct / total

```

```

[16]: all_train_accuracies = [calc_accuracy(trainloader)]
      all_test_accuracies = [calc_accuracy(testloader)]
      for epoch in range(num_epochs): # loop over the dataset multiple times

          running_loss = 0.0
          for i, data in enumerate(trainloader, 0):
              # get the inputs; data is a list of [inputs, labels]
              inputs, labels = data

              # zero the parameter gradients
              optimizer.zero_grad()

              # forward + backward + optimize
              outputs = net(inputs)
              loss = criterion(outputs, labels)
              loss.backward()
              optimizer.step()

              # print statistics
              running_loss += loss.item()
              if i % 2000 == 1999: # print every 2000 mini-batches
                  print('%d, %5d] loss: %.3f' %
                        (epoch + 1, i + 1, running_loss / 2000))
                  running_loss = 0.0

              # Calculate accuracy every 2000 minibatches
              train_accuracy = calc_accuracy(trainloader)
              test_accuracy = calc_accuracy(testloader)
              all_train_accuracies.append(train_accuracy)
              all_test_accuracies.append(test_accuracy)

          print('END OF EPOCH ', epoch + 1, ': train accuracy = ', train_accuracy, ' /
          ↪ test accuracy = ', test_accuracy)

      print('Finished Training')

```

```

[1, 2000] loss: 2.046
[1, 4000] loss: 1.916
[1, 6000] loss: 1.867
[1, 8000] loss: 1.843
[1, 10000] loss: 1.836
[1, 12000] loss: 1.826
END OF EPOCH 1 : train accuracy = 38.652 // test accuracy = 38.15
[2, 2000] loss: 1.798
[2, 4000] loss: 1.788

```

```

[2, 6000] loss: 1.789
[2, 8000] loss: 1.799
[2, 10000] loss: 1.789
[2, 12000] loss: 1.775
END OF EPOCH 2 : train accuracy = 40.252 // test accuracy = 40.04
[3, 2000] loss: 1.776
[3, 4000] loss: 1.759
[3, 6000] loss: 1.747
[3, 8000] loss: 1.770
[3, 10000] loss: 1.753
[3, 12000] loss: 1.756
END OF EPOCH 3 : train accuracy = 41.0 // test accuracy = 40.03
[4, 2000] loss: 1.732
[4, 4000] loss: 1.778
[4, 6000] loss: 1.745
[4, 8000] loss: 1.725
[4, 10000] loss: 1.736
[4, 12000] loss: 1.737
END OF EPOCH 4 : train accuracy = 41.628 // test accuracy = 40.49
[5, 2000] loss: 1.722
[5, 4000] loss: 1.733
[5, 6000] loss: 1.720
[5, 8000] loss: 1.736
[5, 10000] loss: 1.704
[5, 12000] loss: 1.741
END OF EPOCH 5 : train accuracy = 42.056 // test accuracy = 40.7
[6, 2000] loss: 1.710
[6, 4000] loss: 1.720
[6, 6000] loss: 1.727
[6, 8000] loss: 1.728
[6, 10000] loss: 1.704
[6, 12000] loss: 1.717
END OF EPOCH 6 : train accuracy = 42.536 // test accuracy = 40.66
[7, 2000] loss: 1.719
[7, 4000] loss: 1.714
[7, 6000] loss: 1.694
[7, 8000] loss: 1.710
[7, 10000] loss: 1.713
[7, 12000] loss: 1.708
END OF EPOCH 7 : train accuracy = 42.5 // test accuracy = 41.29
[8, 2000] loss: 1.707
[8, 4000] loss: 1.692
[8, 6000] loss: 1.715
[8, 8000] loss: 1.683
[8, 10000] loss: 1.702
[8, 12000] loss: 1.716
END OF EPOCH 8 : train accuracy = 42.544 // test accuracy = 40.82
[9, 2000] loss: 1.708

```

```

[9, 4000] loss: 1.692
[9, 6000] loss: 1.711
[9, 8000] loss: 1.700
[9, 10000] loss: 1.689
[9, 12000] loss: 1.670
END OF EPOCH 9 : train accuracy = 42.868 // test accuracy = 41.13
[10, 2000] loss: 1.695
[10, 4000] loss: 1.674
[10, 6000] loss: 1.703
[10, 8000] loss: 1.694
[10, 10000] loss: 1.687
[10, 12000] loss: 1.692
END OF EPOCH 10 : train accuracy = 43.078 // test accuracy = 40.95
[11, 2000] loss: 1.697
[11, 4000] loss: 1.700
[11, 6000] loss: 1.674
[11, 8000] loss: 1.681
[11, 10000] loss: 1.677
[11, 12000] loss: 1.685
END OF EPOCH 11 : train accuracy = 43.182 // test accuracy = 41.46
[12, 2000] loss: 1.694
[12, 4000] loss: 1.674
[12, 6000] loss: 1.688
[12, 8000] loss: 1.684
[12, 10000] loss: 1.675
[12, 12000] loss: 1.676
END OF EPOCH 12 : train accuracy = 43.556 // test accuracy = 41.33
Finished Training

```


## 4 Plot accuracy over time

```

[17]: import matplotlib.pyplot as plt

plt.figure(1)
plt.plot(all_train_accuracies)
plt.plot(all_test_accuracies)
plt.title('Training and Testing Accuracy after each iteration')
plt.xlabel('Iteration Number')
plt.ylabel('Accuracy (%)')
plt.legend(['Training', 'Testing'])
plt.show()

```



[: