

hw4_p3b

June 5, 2019

```
[1]: %matplotlib inline

[2]: import torch
import torchvision
import torchvision.transforms as transforms

[3]: # LOADS CIFAR10 images which are 32 x 32 x 3 RGB images
# iter(trainloader/testloader) are iterables that come in pairs (image, label)

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

1 Choose Hyperparameters

```
[4]: num_epochs = 15 # number of epochs to train for
momentum = 0.7 # momentum for Stochastic Gradient Descent
lr = 0.001 # learning rate (eta) for gradient descent
M = 200 # number of neurons in hidden layer of neural network
```

2 Build Neural Network

```
[5]: import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self, M):
        super(Net, self).__init__()

        # in features --> flattened 32 x 32 x 3 image, out features = labels
        # (0,1,..., 9), use bias
        self.linear1 = nn.Linear(3072, M, bias=True)
        self.linear2 = nn.Linear(M, 10, bias=True)

    def forward(self, x):
        x = x.view(4, -1)
        x = self.linear1(x)
        x = F.relu(x)
        return self.linear2(x)

net = Net(M)
```

```
[ ]:
```

3. Define a Loss function and optimizer

Let's use a Classification Cross-Entropy loss and SGD with momentum.

```
[6]: import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=lr, momentum=momentum)
```

4. Train the network

This is when things start to get interesting. We simply have to loop over our data iterator, and feed the inputs to the network and optimize.

```
[7]: def calc_accuracy(dataloader):
    correct = 0
    total = 0
    with torch.no_grad():
        for data in dataloader:
            images, labels = data
            outputs = net(images)
```

```

        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    return 100.0 * correct / total

```

```

[8]: all_train_accuracies = []
    all_test_accuracies = []
    for epoch in range(num_epochs): # loop over the dataset multiple times

        running_loss = 0.0
        for i, data in enumerate(trainloader, 0):
            # get the inputs; data is a list of [inputs, labels]
            inputs, labels = data

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = net(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # print statistics
            running_loss += loss.item()
            if i % 2000 == 1999: # print every 2000 mini-batches
                print('[%d, %5d] loss: %.3f' %
                      (epoch + 1, i + 1, running_loss / 2000))
                running_loss = 0.0

            train_accuracy = calc_accuracy(trainloader)
            test_accuracy = calc_accuracy(testloader)
            all_train_accuracies.append(train_accuracy)
            all_test_accuracies.append(test_accuracy)

        print('END OF EPOCH ', epoch + 1, ': train accuracy = ', train_accuracy, ' /
        →/ test accuracy = ', test_accuracy)

    print('Finished Training')

```

```

[1, 2000] loss: 1.912
[1, 4000] loss: 1.716
[1, 6000] loss: 1.658
[1, 8000] loss: 1.631
[1, 10000] loss: 1.603

```

```

[1, 12000] loss: 1.584
END OF EPOCH 1 : train accuracy = 46.69 // test accuracy = 45.28
[2, 2000] loss: 1.524
[2, 4000] loss: 1.510
[2, 6000] loss: 1.474
[2, 8000] loss: 1.463
[2, 10000] loss: 1.460
[2, 12000] loss: 1.432
END OF EPOCH 2 : train accuracy = 51.152 // test accuracy = 48.19
[3, 2000] loss: 1.394
[3, 4000] loss: 1.390
[3, 6000] loss: 1.386
[3, 8000] loss: 1.389
[3, 10000] loss: 1.371
[3, 12000] loss: 1.367
END OF EPOCH 3 : train accuracy = 54.648 // test accuracy = 50.4
[4, 2000] loss: 1.300
[4, 4000] loss: 1.310
[4, 6000] loss: 1.318
[4, 8000] loss: 1.313
[4, 10000] loss: 1.325
[4, 12000] loss: 1.321
END OF EPOCH 4 : train accuracy = 57.4 // test accuracy = 51.49
[5, 2000] loss: 1.241
[5, 4000] loss: 1.289
[5, 6000] loss: 1.256
[5, 8000] loss: 1.233
[5, 10000] loss: 1.279
[5, 12000] loss: 1.256
END OF EPOCH 5 : train accuracy = 60.178 // test accuracy = 52.21
[6, 2000] loss: 1.207
[6, 4000] loss: 1.195
[6, 6000] loss: 1.204
[6, 8000] loss: 1.210
[6, 10000] loss: 1.222
[6, 12000] loss: 1.227
END OF EPOCH 6 : train accuracy = 60.5 // test accuracy = 51.89
[7, 2000] loss: 1.130
[7, 4000] loss: 1.149
[7, 6000] loss: 1.170
[7, 8000] loss: 1.193
[7, 10000] loss: 1.164
[7, 12000] loss: 1.191
END OF EPOCH 7 : train accuracy = 61.702 // test accuracy = 51.35
[8, 2000] loss: 1.090
[8, 4000] loss: 1.118
[8, 6000] loss: 1.132
[8, 8000] loss: 1.125

```

```

[8, 10000] loss: 1.133
[8, 12000] loss: 1.158
END OF EPOCH 8 : train accuracy = 64.118 // test accuracy = 52.27
[9, 2000] loss: 1.087
[9, 4000] loss: 1.087
[9, 6000] loss: 1.075
[9, 8000] loss: 1.097
[9, 10000] loss: 1.090
[9, 12000] loss: 1.109
END OF EPOCH 9 : train accuracy = 64.516 // test accuracy = 52.07
[10, 2000] loss: 1.030
[10, 4000] loss: 1.057
[10, 6000] loss: 1.038
[10, 8000] loss: 1.052
[10, 10000] loss: 1.097
[10, 12000] loss: 1.071
END OF EPOCH 10 : train accuracy = 66.502 // test accuracy = 52.73
[11, 2000] loss: 1.008
[11, 4000] loss: 1.016
[11, 6000] loss: 1.006
[11, 8000] loss: 1.033
[11, 10000] loss: 1.032
[11, 12000] loss: 1.057
END OF EPOCH 11 : train accuracy = 68.582 // test accuracy = 52.41
[12, 2000] loss: 0.937
[12, 4000] loss: 0.962
[12, 6000] loss: 0.985
[12, 8000] loss: 1.021
[12, 10000] loss: 0.993
[12, 12000] loss: 1.038
END OF EPOCH 12 : train accuracy = 68.81 // test accuracy = 52.11
[13, 2000] loss: 0.926
[13, 4000] loss: 0.943
[13, 6000] loss: 0.963
[13, 8000] loss: 0.973
[13, 10000] loss: 0.969
[13, 12000] loss: 1.007
END OF EPOCH 13 : train accuracy = 69.066 // test accuracy = 51.63
[14, 2000] loss: 0.906
[14, 4000] loss: 0.904
[14, 6000] loss: 0.929
[14, 8000] loss: 0.962
[14, 10000] loss: 0.952
[14, 12000] loss: 0.952
END OF EPOCH 14 : train accuracy = 71.012 // test accuracy = 51.68
[15, 2000] loss: 0.867
[15, 4000] loss: 0.894
[15, 6000] loss: 0.905

```

```
[15, 8000] loss: 0.915
[15, 10000] loss: 0.932
[15, 12000] loss: 0.921
END OF EPOCH 15 : train accuracy = 69.858 // test accuracy = 50.2
Finished Training
```

5 Plot accuracy over time

```
[9]: import matplotlib.pyplot as plt

plt.figure(1)
plt.plot(all_train_accuracies)
plt.plot(all_test_accuracies)
plt.title('Training and Testing Accuracy after each iteration')
plt.xlabel('Iteration Number')
plt.ylabel('Accuracy (%)')
plt.legend(['Training', 'Testing'])
plt.show()
```

