

Linear Regression: Model and Algorithms (part 2)

CSE 446

Most slides by Emily Fox
Presented by Anna Karlin
April 8, 2019



XKCD

Linear regression: a supervised learning problem

Goal: to predict some output from some inputs using labelled examples. Example: house sales price from square footage

Supervised learning: Problem of learning a function that maps inputs to outputs based on labelled examples.

Regression: When the labels are real numbers

Linear regression: a supervised learning problem

Goal: to predict some output from some inputs/features. Example: house sales price from square footage

Step 1: Define set up and get data

- a model for how the output y depends on the inputs \mathbf{x} .
- We assumed that y is a linear function of features + noise.

$$y = \mathbf{w}^T \mathbf{x} + \varepsilon$$

- A training set (labelled examples): $\{(\mathbf{x}_i, y_i)\}_{i=1..n}$

Linear regression: a supervised learning problem

Goal: to predict some output from some inputs/features.

Step 1: Define set up ($y = \mathbf{w}^T \mathbf{x} + \varepsilon$) and get data $\{(\mathbf{x}_i, y_i)\}_{i=1..n}$

Step 2: find the parameters \mathbf{w} that minimize the “loss/cost” on the training set.

- Our **loss function** was residual sum of squares (RSS)
- Find $\hat{\mathbf{w}}$ that minimizes $\text{RSS} = \sum_{i=1}^n (y_i - \sum_{j=1}^d \mathbf{w}_j x_i[j])^2 = (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w})$
- Found solution by solving for gradient of $(\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) = \mathbf{0}$

Solution: $\mathbf{X}^T \mathbf{X} \hat{\mathbf{w}} = \mathbf{X}^T \mathbf{y}$ If $\mathbf{X}^T \mathbf{X}$ is invertible, could write $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

Linear regression: a supervised learning problem

Goal: to predict some output from some inputs/features.

Step 1: Define set up ($y = \mathbf{w}^T \mathbf{x} + \epsilon$) and get data $\{(\mathbf{x}_i, y_i)\}_{i=1..n}$

Step 2: find the parameters $\hat{\mathbf{w}}$ that minimizes $RSS = (\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w})$

Step 3: Use $\hat{\mathbf{w}}$ to make predictions.

Given \mathbf{x} , predict output: $\hat{\mathbf{w}}^T \mathbf{x}$

Plan for today:

- Gradient descent
- Handling an intercept
- More features/more complex models
- How well does it work?

Linear regression: a supervised learning problem

Goal: to predict some output from some inputs/features. Example: house sales price from square footage

Step 1: Define set up and get data

- a model for how the output y depends on the inputs x .
- We assumed that y is a linear function of features + noise.

$$y = \mathbf{w}^T \mathbf{x} + \varepsilon$$

- A training set (labelled examples): $\{(\mathbf{x}_i, y_i)\}_{i=1..n}$

Step 2: find the parameters \mathbf{w} that minimize the “loss/cost” on the training set.

- Our **loss function** was residual sum of squares (RSS)
- Find $\hat{\mathbf{w}}$ that minimizes $\text{RSS} = \sum_{i=1}^n (y_i - \sum_{j=1}^d \mathbf{x}_i^T \mathbf{w}_j)^2 = (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w})$
- Found solution by solving for gradient of $(\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) = \mathbf{0}$

$$\text{Solution: } \mathbf{X}^T \mathbf{X} \hat{\mathbf{w}} = \mathbf{X}^T \mathbf{y}$$

If $\mathbf{X}^T \mathbf{X}$ is invertible, could write $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

Step 3: Use $\hat{\mathbf{w}}$ to make predictions, Given \mathbf{x} , predict output: $\mathbf{w}^T \mathbf{x}$

Fitting the linear regression model


Gradient descent

Gradient Descent – univariate case

- Repeatedly move in direction that reduces the value of the function.

Gradient Descent – multivariate case


Gradient descent for linear regression: repeatedly move in direction of negative gradient



while not converged

$$w^{(t+1)} \leftarrow w^{(t)} - \eta \underbrace{\nabla \text{RSS}(w^{(t)})}_{-2X^T(y - Xw^{(t)})}$$


Interpreting elementwise



Update to j^{th} feature weight:

$$w_j^{(t+1)} \leftarrow w_j^{(t)} + 2\eta \sum_{i=1}^N x_i[j] (y_i - \hat{y}_i(\mathbf{w}^{(t)}))$$

Summary of gradient descent for multiple regression



init $\mathbf{w}^{(1)} = \mathbf{0}$ (or randomly, or smartly), $t = 1$

while $|| \nabla \text{RSS}(\mathbf{w}^{(t)}) || > \epsilon$

for $j = 1, \dots, d$

partial[j] = $-2 \sum_{i=1}^n x_i[j] (y_i - \hat{y}_i(\mathbf{w}^{(t)}))$

$w_j^{(t+1)} \leftarrow w_j^{(t)} - \eta \text{partial}[j]$

$t \leftarrow t + 1$

Adding an intercept – “demeaning”

Once we have a fitted function

- We use it to predict the sales price for new houses, by plugging in square footage, number of bathrooms, etc for the new house \mathbf{x} whose sales price we want to predict.
- Prediction is:

What if we want to allow for an intercept?

Assume that $y = \mathbf{w}^T \mathbf{x} + b + \varepsilon$

Find $\hat{\mathbf{w}}, b$ that minimize $\text{RSS} = \sum_{i=1}^n (y_i - \sum_{j=1}^d w_j x_i[j] - b)^2$

$$= (\mathbf{y} - \mathbf{X}\mathbf{w} - b\mathbf{1})^T (\mathbf{y} - \mathbf{X}\mathbf{w} - b\mathbf{1})$$

Handling an intercept (constant term)

Assume that $y = \mathbf{w}^T \mathbf{x} + b + \varepsilon$

Find $\hat{\mathbf{w}}, b$ that minimize
$$\text{RSS} = \sum_{i=1}^n (y_i - \sum_{j=1}^d w_j x_i[j] - b)^2$$
$$= (\mathbf{y} - \mathbf{X}\mathbf{w} - b\mathbf{1})^T (\mathbf{y} - \mathbf{X}\mathbf{w} - b\mathbf{1})$$

Two step approach:

1. Show that if $\frac{1}{n} \sum_i \mathbf{x}_i = \mathbf{0}$ (*) then solution is simple.
2. Show how to transform, aka “demean” any linear regression problem so that (*) holds.

1. Show that if $\frac{1}{n} \sum_i \mathbf{x}_i = \mathbf{0}$ (*) then solution is simple.

Same as saying that $\mathbf{X}^T \mathbf{1} = \mathbf{0}$.

Find $\hat{\mathbf{w}}, b$ that minimize

$$\begin{aligned} \text{RSS} &= \sum_{i=1}^n (y_i - \sum_{j=1}^d w_j x_i[j] - b)^2 \\ &= (\mathbf{y} - \mathbf{X}\mathbf{w} - b\mathbf{1})^T (\mathbf{y} - \mathbf{X}\mathbf{w} - b\mathbf{1}) \end{aligned}$$

$$\text{partial}[w_j] = -2 \sum_{i=1}^n x_i[j] (y_i - \mathbf{x}_i^T \mathbf{w} - b)$$


$$\text{partial}[b] = -2 \sum_{i=1}^n (y_i - \mathbf{x}_i^T \mathbf{w} - b)$$

2. Show how to transform, aka “demean” any linear regression problem so that (*) holds.


$$\frac{1}{n} \sum_i \mathbf{x}_i = \mathbf{0} \quad (*)$$

More features, more complex models

What about a quadratic function?



Even higher order polynomial




Polynomial regression (single input)

Goal: to predict some output from some inputs/features.

Step 1: Assume that y (sales price) is a polynomial function of feature (square footage)+ noise.

$$y_i = \sum_{j=0}^p w_j x_i^j + \epsilon \quad \text{A training set (labelled examples):}$$



Polynomial regression

Goal: to predict some output from some inputs/features.

Step 1: $y_i = \sum_{j=0}^p w_j x_i^j + \varepsilon$ A training set (labelled examples):

Step 2: find params \mathbf{w} that minimize the “loss/cost” on training set $\{(x_i, y_i)\}_{i=1..n}$

- **Loss function** is residual sum of squares (RSS)
- Find $\hat{\mathbf{w}}$ that minimizes
$$\text{RSS} = \sum_{i=1}^n (y_i - \sum_{j=0}^{p-1} w_j x_i^j)^2$$

Polynomial regression

Goal: to predict some output from some inputs/features.

Step 1: $y_i = \sum_{j=0}^p w_j x_i^j + \varepsilon$ A training set (labelled examples):

Step 2: find params w that minimize the “loss/cost” on training set $\{(x_i, y_i)\}_{i=1..n}$

- Find \hat{w} that minimizes $RSS = \sum_{i=1}^n (y_i - \sum_{j=0}^{p-1} w_j x_i^j)^2$
- **Just as easy to solve!** Just think of x_i^j as one of p features associated with the i^{th} observation.
- Instead of single input x_i , define features $\mathbf{h}(x) = (1, x, x^2, \dots, x^p)$

$$\begin{aligned}\mathbf{h}(x_i) &= (h_0(x_i), h_1(x_i), h_2(x_i), h_3(x_i), h_4(x_i), h_5(x_i)) \\ &= (1, x_i, x_i^2, x_i^3, x_i^4, x_i^5)\end{aligned}$$

Polynomial regression

Step 1: $y_i = \sum_{j=0}^p \mathbf{w}_j x_i^j + \varepsilon$

Step 2: find the parameters \mathbf{w} that minimize the “loss/cost” on the training set.

- Find $\hat{\mathbf{w}}$ that minimizes $\text{RSS} = \sum_{i=1}^n (y_i - \sum_{j=0}^p \mathbf{w}_j x_i^j)^2 = (\mathbf{y} - \mathbf{H}\mathbf{w})^\top (\mathbf{y} - \mathbf{H}\mathbf{w})$
- Find solution by solving for gradient of $(\mathbf{y} - \mathbf{H}\mathbf{w})^\top (\mathbf{y} - \mathbf{H}\mathbf{w}) = 0$

Solution: $\mathbf{H}^\top \mathbf{H} \hat{\mathbf{w}} = \mathbf{H}^\top \mathbf{y}$

Polynomial regression

Step 1: $y_i = \sum_{j=0}^p w_j x_i^j + \varepsilon$

Step 2: find the parameters \mathbf{w} that minimize the “loss/cost” on the training set.

- Find $\hat{\mathbf{w}}$ that minimizes $\text{RSS} = \sum_{i=1}^n (y_i - \sum_{j=0}^p w_j x_i^j)^2 = (\mathbf{y} - \mathbf{H}\mathbf{w})^T(\mathbf{y} - \mathbf{H}\mathbf{w})$
- Find solution by solving for gradient of $(\mathbf{y} - \mathbf{H}\mathbf{w})^T(\mathbf{y} - \mathbf{H}\mathbf{w}) = 0$

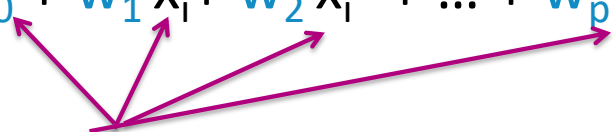
Solution: $\mathbf{H}^T \mathbf{H} \hat{\mathbf{w}} = \mathbf{H}^T \mathbf{y}$

Step 3: Use $\hat{\mathbf{w}} = (\hat{w}_0, \hat{w}_1, \dots, \hat{w}_p)$ to make predictions. Given x , let $\mathbf{h}(x) = (1, x, x^2, \dots, x^p)$ and predict output:

$$\mathbf{h}(x)^T \hat{\mathbf{w}} = \hat{w}_0 + \hat{w}_1 x + \dots + \hat{w}_p x^p$$

Polynomial regression

Model:

$$y_i = w_0 + w_1 x_i + w_2 x_i^2 + \dots + w_p x_i^p + \epsilon_i$$


treat transformed inputs as different features

feature 1 = 1 (constant)

feature 2 = x

feature 3 = x^2

...

feature $p+1 = x^p$

parameter 1 = w_0

parameter 2 = w_1

parameter 3 = w_2

...

parameter $p+1 = w_p$

Why might we want to use polynomial regression?

- Taylor Series!

More generally

- Start with set of inputs for each observation $\mathbf{x} = (x[1], x[2], \dots, x[d])$ and training set: $\{(\mathbf{x}_i, y_i)\}_{i=1..n}$
- Define feature map that transforms each input vector \mathbf{x}_i to higher dimensional feature vector $h(\mathbf{x}_i)$.

Example: $x_i[1]$ $x_i[2]$ $x_i[3]$

$h_1(\mathbf{x})$	$h_2(\mathbf{x}_i)$	$h_3(\mathbf{x}_i)$	$h_4(\mathbf{x}_i),$	$h_5(\mathbf{x}_i)$	$h_6(\mathbf{x}_i)$	$h_7(\mathbf{x}_i)$
1	$x_i[1]$	$x_i[1]^2$	$x_i[1]x_i[2]$	$x_i[2]$	$x_i[2]^2$	$\cos(\pi x_i[3]/6)$

General notation

Output: y  *scalar*

Inputs: $\mathbf{x} = (x[1], x[2], \dots, x[d])$

 *d-dim vector*

Notational conventions:

\mathbf{x}_i = input of i^{th} data point (*vector*)

$x_i[j]$ = j^{th} input of i^{th} data point (*scalar*)

$\mathbf{h}(\mathbf{x}) = (h_1(\mathbf{x}), h_2(\mathbf{x}), \dots, h_p(\mathbf{x}))$ feature map applied to input \mathbf{x} (*vector*)

$h_j(\mathbf{x})$ = j^{th} feature associated with input \mathbf{x} (*scalar*) (j^{th} basis function)

H = n by p matrix whose i^{th} row is $\mathbf{h}(\mathbf{x}_i)$

To fit these more general functions

- Start with input features $\mathbf{x} = (x[1], x[2], \dots, x[d])$
and training set: $\{(\mathbf{x}_i, y_i)\}_{i=1..n}$
- Define feature map that transforms each \mathbf{x}_i to higher dimensional feature vector $\mathbf{h}(\mathbf{x}_i)$.
- Model: $y_i = \sum_{j=1}^p \mathbf{w}_j \mathbf{h}_j(\mathbf{x}_i) + \varepsilon_i$
- Find $\hat{\mathbf{w}}$ that minimizes
$$\text{RSS} = \sum_{i=1}^n (y_i - \sum_{j=1}^p \mathbf{w}_j \mathbf{h}_j(\mathbf{x}_i))^2$$
$$= (\mathbf{y} - \mathbf{H}\mathbf{w})^T (\mathbf{y} - \mathbf{H}\mathbf{w})$$
- Solution: $\mathbf{H}^T \mathbf{H} \hat{\mathbf{w}} = \mathbf{H}^T \mathbf{y}$

Recap of concepts

What you can do now...

- Describe linear regression (and feature maps)
- Write a regression model using multiple inputs or features thereof.
- Calculate a goodness-of-fit metric (e.g., RSS)
- Estimate model parameters of a general multiple regression model to minimize RSS:
 - In closed form
 - Using an iterative gradient descent algorithm
- Interpret the coefficients of a non-featurized multiple regression fit
- Exploit the estimated model to form predictions