

# Warm up

```
# generate some nonsense data for an example
X = np.random.randn(n,d)
y = np.random.randn(n)
```

```
# generate the random features
G = np.random.randn(p, d)*np.sqrt(.1)
b = np.random.rand(p)*2*np.pi
```

```
# construct HTH
HTH = np.zeros((p,p))
HTy = np.zeros(p)
for i in range(n):
    hi = np.dot(X[i,:], G.T)+b
    HTH += np.outer(hi, hi)
    HTy += y[i]*hi
    if i % 1000==0: print(i)
```

```
H = np.dot(X, G.T) + b.T
HTH = np.dot(H.T, H)
HTy = np.dot(H.T, y)
```

```
# construct HTH
HTH = np.zeros((p,p))
HTy = np.zeros(p)
block = p
for i in range(int(np.ceil(n/block))+1):
    Hi = np.dot(X[i*block:min(n,(i+1)*block),:], G.T)+b
    HTH += np.dot(Hi.T, Hi)
    HTy += np.dot(Hi.T, y[i*block:min(n,(i+1)*block)])
```

```
w = np.linalg.solve(HTH + lam*np.eye(p), HTy)
```

1 float in NumPy = 8 bytes

$10^6 \approx 2^{20}$  bytes = 1 MB

$10^9 \approx 2^{30}$  bytes = 1 GB

For each block compute the memory required in terms of  $n$ ,  $p$ ,  $d$ .

If  $d \ll p \ll n$ , what is the most memory efficient program (blue, green, red)?

If you have unlimited memory, what do you think is the fastest program?



# Regularization

Machine Learning – CSE546

Kevin Jamieson

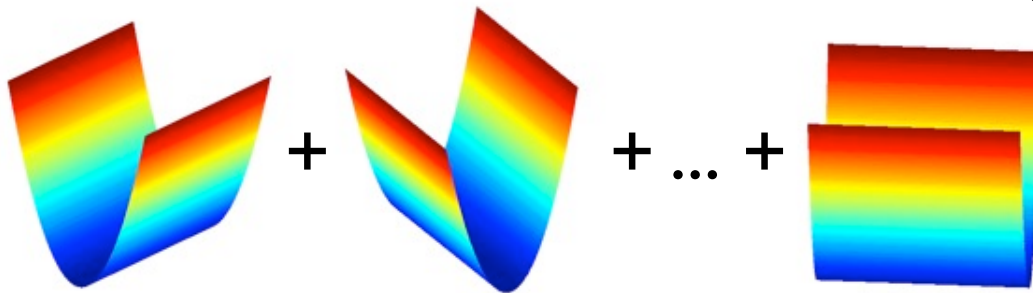
University of Washington

April 15, 2019

# Ridge Regression

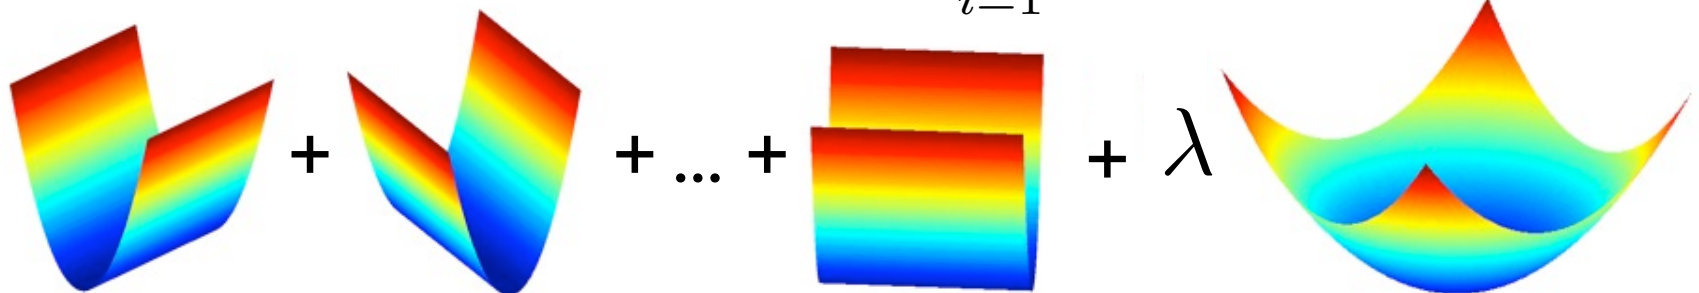
- Old Least squares objective:

$$\hat{w}_{LS} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2$$



- Ridge Regression objective:

$$\hat{w}_{ridge} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2$$



# Minimizing the Ridge Regression Objective

$$\hat{w}_{ridge} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2$$

$$0 = \nabla_w \left( \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2 \right)$$

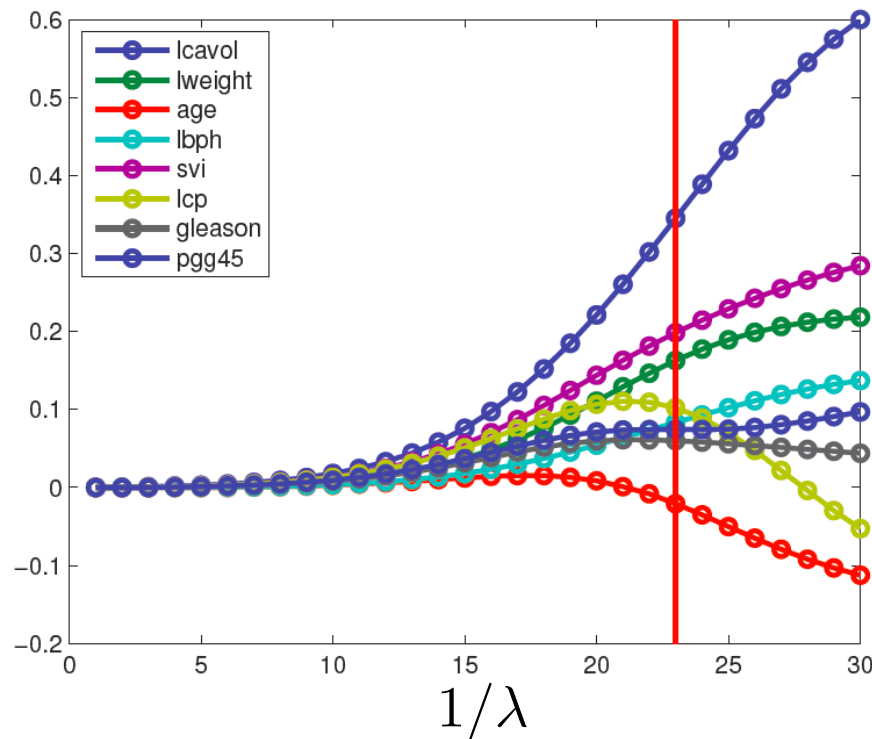
$$= - \sum_{i=1}^n 2x_i (y_i - x_i^T w) + 2\lambda w$$

$$= 2 \left( \sum_{i=1}^n x_i y_i \right) + 2 \left( \sum_{i=1}^n x_i x_i^T + \lambda I \right) w$$

$$\begin{aligned} \hat{w}_{ridge} &= \left( \sum_{i=1}^n x_i x_i^T + \lambda I \right)^{-1} \left( \sum_{i=1}^n x_i y_i \right) \\ &= (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y} \end{aligned}$$

# Ridge Coefficient Path

$\mathbf{X}^T \mathbf{X}$  in general



From  
Kevin Murphy  
textbook

- Typical approach: select  $\lambda$  using cross validation, up next

# Bias-Variance Properties

$$\hat{w}_{ridge} = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$$

- **Assume:**  $\mathbf{X}^T \mathbf{X} = nI$  **and**  $\mathbf{y} = \mathbf{X}w + \epsilon$   $\epsilon \sim \mathcal{N}(0, \sigma^2 I)$

If  $x \in \mathbb{R}^d$  and  $Y \sim \mathcal{N}(x^T w, \sigma^2)$ , what is  $\mathbb{E}_{Y|x, \text{train}}[(Y - x^T \hat{w}_{ridge})^2 | X = x]$ ?

$$\begin{aligned} & \mathbb{E}_{Y|X, \mathcal{D}}[(Y - x^T \hat{w}_{ridge})^2 | X = x] \\ &= \mathbb{E}_{Y|X}[(Y - \mathbb{E}_{Y|X}[Y|X = x])^2 | X = x] + \mathbb{E}_{\mathcal{D}}[(\mathbb{E}_{Y|X}[Y|X = x] - x^T \hat{w}_{ridge})^2] \\ &= \mathbb{E}_{Y|X}[(Y - x^T w)^2 | X = x] + \mathbb{E}_{\mathcal{D}}[(x^T w - x^T \hat{w}_{ridge})^2] \\ &= \sigma^2 + (x^T w - \mathbb{E}_{\mathcal{D}}[x^T \hat{w}_{ridge}])^2 + \mathbb{E}_{\mathcal{D}}[(\mathbb{E}_{\mathcal{D}}[x^T \hat{w}_{ridge}] - x^T \hat{w}_{ridge})^2] \\ &= \sigma^2 + \frac{\lambda^2}{(n + \lambda)^2} (w^T x)^2 + \frac{d\sigma^2 n}{(n + \lambda)^2} \|x\|_2^2 \end{aligned}$$

Irreduc. Error      Bias-squared      Variance      (verify at home)

# Ridge Regression: Effect of Regularization

$$\mathcal{D} \stackrel{i.i.d.}{\sim} P_{XY} \quad \hat{w}_{\mathcal{D},ridge}^{(\lambda)} = \arg \min_w \frac{1}{|\mathcal{D}|} \sum_{(x_i, y_i) \in \mathcal{D}} (y_i - x_i^T w)^2 + \lambda \|w\|_2^2$$

**TRAIN error:**

$$\frac{1}{|\mathcal{D}|} \sum_{(x_i, y_i) \in \mathcal{D}} (y_i - x_i^T \hat{w}_{\mathcal{D},ridge}^{(\lambda)})^2$$

# Ridge Regression: Effect of Regularization

$$\mathcal{D} \stackrel{i.i.d.}{\sim} P_{XY} \quad \hat{w}_{\mathcal{D},ridge}^{(\lambda)} = \arg \min_w \frac{1}{|\mathcal{D}|} \sum_{(x_i, y_i) \in \mathcal{D}} (y_i - x_i^T w)^2 + \lambda \|w\|_2^2$$

**TRAIN error:**

$$\frac{1}{|\mathcal{D}|} \sum_{(x_i, y_i) \in \mathcal{D}} (y_i - x_i^T \hat{w}_{\mathcal{D},ridge}^{(\lambda)})^2$$

**TEST error:**

$$\mathcal{T} \stackrel{i.i.d.}{\sim} P_{XY}$$

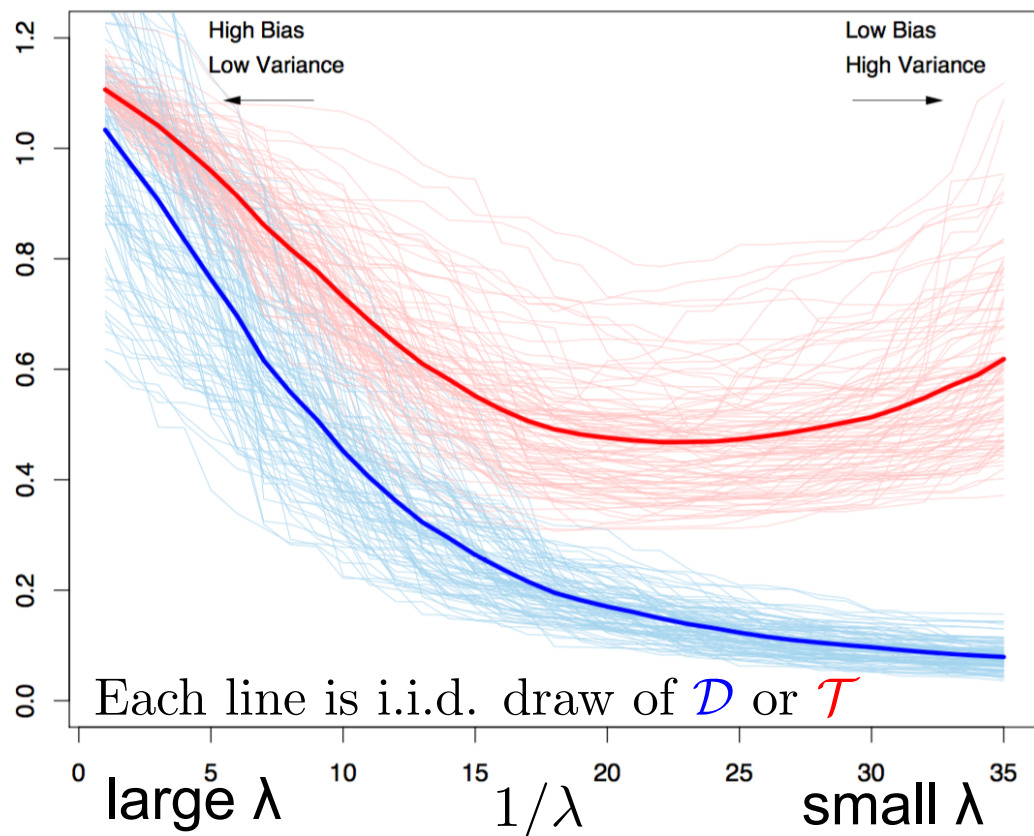
$$\frac{1}{|\mathcal{T}|} \sum_{(x_i, y_i) \in \mathcal{T}} (y_i - x_i^T \hat{w}_{\mathcal{D},ridge}^{(\lambda)})^2$$

Important:  $\mathcal{D} \cap \mathcal{T} = \emptyset$



# Ridge Regression: Effect of Regularization

$$\mathcal{D} \stackrel{i.i.d.}{\sim} P_{XY} \quad \hat{w}_{\mathcal{D},ridge}^{(\lambda)} = \arg \min_w \frac{1}{|\mathcal{D}|} \sum_{(x_i, y_i) \in \mathcal{D}} (y_i - x_i^T w)^2 + \lambda \|w\|_2^2$$



**TRAIN error:**

$$\frac{1}{|\mathcal{D}|} \sum_{(x_i, y_i) \in \mathcal{D}} (y_i - x_i^T \hat{w}_{\mathcal{D},ridge}^{(\lambda)})^2$$

**TEST error:**

$$\mathcal{T} \stackrel{i.i.d.}{\sim} P_{XY}$$

$$\frac{1}{|\mathcal{T}|} \sum_{(x_i, y_i) \in \mathcal{T}} (y_i - x_i^T \hat{w}_{\mathcal{D},ridge}^{(\lambda)})^2$$

Important:  $\mathcal{D} \cap \mathcal{T} = \emptyset$



# Cross-Validation

Machine Learning – CSE546

Kevin Jamieson

University of Washington

October 9, 2016

# How... How... How???????

- *How do we pick the regularization constant  $\lambda$ ...*
- *How do we pick the number of basis functions...*
- We could use the test data, but...

# How... How... How???????

- *How do we pick the regularization constant  $\lambda$ ...*
- *How do we pick the number of basis functions...*
- We could use the test data, but...
- Never ever ever ever ever ever ever ever  
ever ever ever ever ever ever ever ever  
ever ever ever ever ever ever ever ever  
train on the test data

# (LOO) Leave-one-out cross validation

- Consider a **validation set with 1 example**:
  - $D$  – training data
  - $D \setminus j$  – training data with  $j$ th data point  $(\mathbf{x}_j, \mathbf{y}_j)$  moved to validation set
- **Learn classifier  $f_{D \setminus j}$  with  $D \setminus j$  dataset**
- **Estimate true error** as squared error on predicting  $\mathbf{y}_j$ :
  - Unbiased estimate of error<sub>true</sub>( $f_{D \setminus j}$ )!

□

# (LOO) Leave-one-out cross validation

- Consider a **validation set with 1 example**:
  - $D$  – training data
  - $D \setminus j$  – training data with  $j$ th data point  $(\mathbf{x}_j, \mathbf{y}_j)$  moved to validation set
- **Learn classifier  $f_{D \setminus j}$  with  $D \setminus j$  dataset**
- **Estimate true error** as squared error on predicting  $\mathbf{y}_j$ :
  - Unbiased estimate of  $\text{error}_{\text{true}}(f_{D \setminus j})$ !
- **LOO cross validation**: Average over all data points  $j$ :
  - For each data point you leave out, learn a new classifier  $f_{D \setminus j}$
  - Estimate error as:

$$\text{error}_{LOO} = \frac{1}{n} \sum_{j=1}^n (y_j - f_{D \setminus j}(x_j))^2$$

# LOO cross validation is (almost) unbiased estimate of true error of $h_D$ !

- When computing **LOOCV error**, we only use  **$N-1$  data points**
  - So it's not estimate of true error of learning with  $N$  data points
  - Usually pessimistic, though – learning with less data typically gives worse answer
- **LOO is almost unbiased! Use LOO error for model selection!!!**
  - **E.g., picking  $\lambda$**

# Computational cost of LOO

- Suppose you have 100,000 data points
- You implemented a great version of your learning algorithm
  - Learns in only 1 second
- Computing LOO will take about 1 day!!!
  -



# Use $k$ -fold cross validation

- Randomly divide training data into  $k$  equal parts

- $D_1, \dots, D_k$

- For each  $i$

- Learn classifier  $f_{D \setminus D_i}$  using data point not in  $D_i$

- Estimate error of  $f_{D \setminus D_i}$  on validation set  $D_i$ :

$$\text{error}_{D_i} = \frac{1}{|D_i|} \sum_{(x_j, y_j) \in D_i} (y_j - f_{D \setminus D_i}(x_j))^2$$

1	2	3	4	5
Train	Train	Validation	Train	Train

# Use $k$ -fold cross validation

- Randomly divide training data into  $k$  equal parts

- $D_1, \dots, D_k$

- For each  $i$

- Learn classifier  $f_{D \setminus D_i}$  using data point not in  $D_i$

- Estimate error of  $f_{D \setminus D_i}$  on validation set  $D_i$ :

$$\text{error}_{\mathcal{D}_i} = \frac{1}{|\mathcal{D}_i|} \sum_{(x_j, y_j) \in \mathcal{D}_i} (y_j - f_{\mathcal{D} \setminus \mathcal{D}_i}(x_j))^2$$

- $k$ -fold cross validation error is average over data splits:

$$\text{error}_{k\text{-fold}} = \frac{1}{k} \sum_{i=1}^k \text{error}_{\mathcal{D}_i}$$

- $k$ -fold cross validation properties:

- Much faster to compute than LOO

- More (pessimistically) biased – using much less data, only  $n(k-1)/k$

- Usually,  $k = 10$

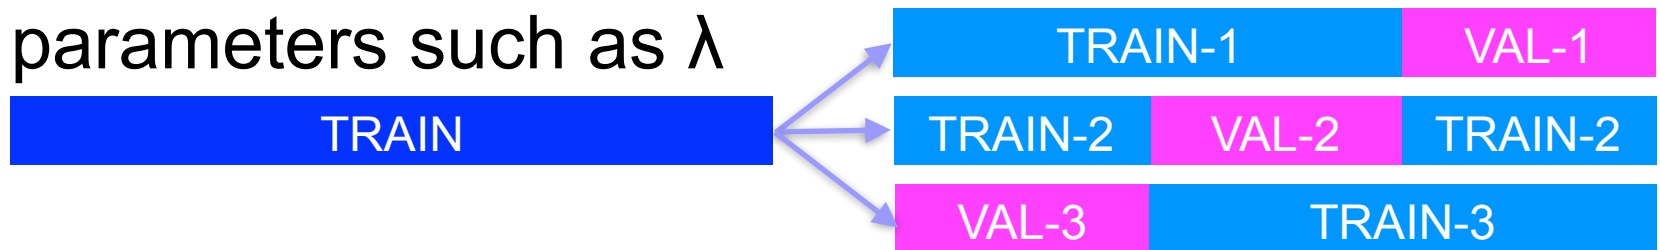
1	2	3	4	5
Train	Train	Validation	Train	Train

# Recap

- Given a dataset, begin by splitting into



- Model selection:** Use k-fold cross-validation on **TRAIN** to train predictor and choose magic parameters such as  $\lambda$



- Model assessment:** Use **TEST** to assess the accuracy of the model you output
  - Never ever ever ever ever train or choose parameters based on the test data

# Example

- Given 10,000-dimensional data and  $n$  examples, we pick a subset of 50 dimensions that have the highest correlation with labels in the training set:

50 indices  $j$  that have largest  $\frac{|\sum_{i=1}^n x_{i,j} y_i|}{\sqrt{\sum_{i=1}^n x_{i,j}^2}}$

- After picking our 50 features, we then use CV to train ridge regression with regularization  $\lambda$
- What's wrong with this procedure?

# Recap

- Learning is...
  - Collect some data
    - E.g., housing info and sale price
  - Randomly split dataset into TRAIN, VAL, and TEST
    - E.g., 80%, 10%, and 10%, respectively
  - Choose a hypothesis class or model
    - E.g., linear with non-linear transformations
  - Choose a loss function
    - E.g., least squares with ridge regression penalty on TRAIN
  - Choose an optimization procedure
    - E.g., set derivative to zero to obtain estimator, cross-validation on VAL to pick num. features and amount of regularization
  - Justifying the accuracy of the estimate
    - E.g., report TEST error



# Simple Variable Selection LASSO: Sparse Regression

Machine Learning – CSE546

Kevin Jamieson

University of Washington

October 9, 2016

# Sparsity

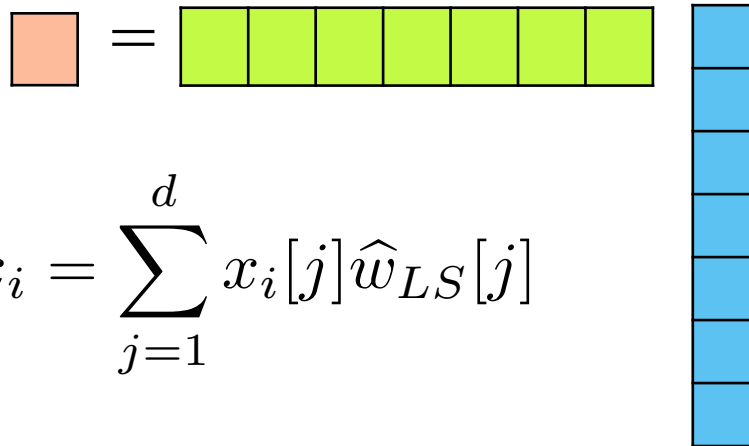
$$\hat{w}_{LS} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2$$

- Vector  $\mathbf{w}$  is sparse, if many entries are zero

# Sparsity

$$\hat{w}_{LS} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2$$

- Vector  $\mathbf{w}$  is sparse, if many entries are zero
- **Efficiency:** If  $\text{size}(\mathbf{w}) = 100$  Billion, each prediction is expensive:
  - If  $\mathbf{w}$  is sparse, prediction computation only depends on number of non-zeros



$$\hat{y}_i = \hat{w}_{LS}^T x_i = \sum_{j=1}^d x_i[j] \hat{w}_{LS}[j]$$



# Sparsity

$$\hat{w}_{LS} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2$$

- Vector  $w$  is sparse, if many entries are zero

- **Interpretability:** What are the relevant dimension to make a prediction?



- How do we find “best” subset among all possible?

Lot size	Dishwasher
Single Family	Garbage disposal
Year built	Microwave
Last sold price	Range / Oven
Last sale price/sqft	Refrigerator
Finished sqft	Washer
Unfinished sqft	Dryer
Finished basement sqft	Laundry location
# floors	Heating type
Flooring types	Jetted Tub
Parking type	Deck
Parking amount	Fenced Yard
Cooling	Lawn
Heating	Garden
Exterior materials	Sprinkler System
Roof type	
Structure style	

# Finding best subset: **Exhaustive**



- Try all subsets of size 1, 2, 3, ... and one that minimizes validation error
- Problem?

# Finding best subset: **Greedy**



## **Forward stepwise:**

Starting from simple model and iteratively add features most useful to fit

## **Backward stepwise:**

Start with full model and iteratively remove features least useful to fit

## **Combining forward and backward steps:**

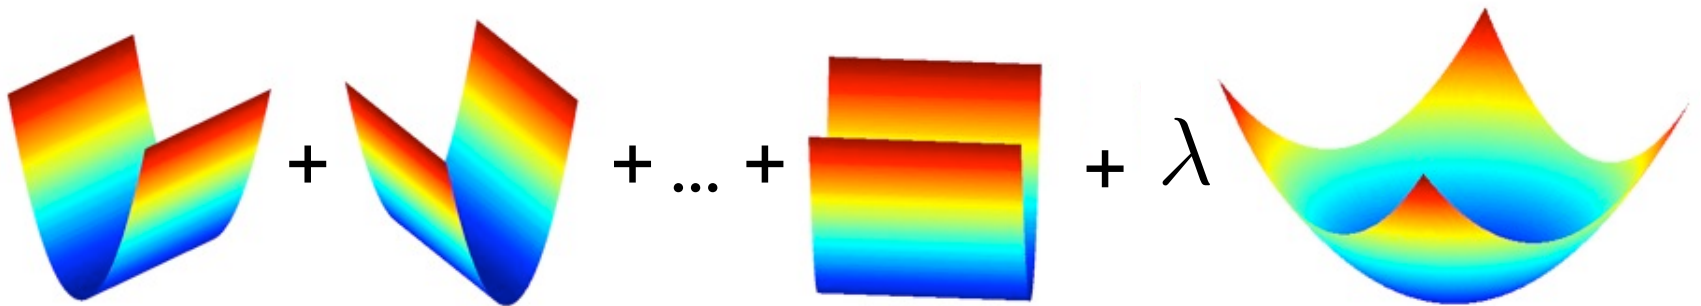
In forward algorithm, insert steps to remove features no longer as important

*Lots of other variants, too.*

# Finding best subset: Regularize

Ridge regression makes coefficients small

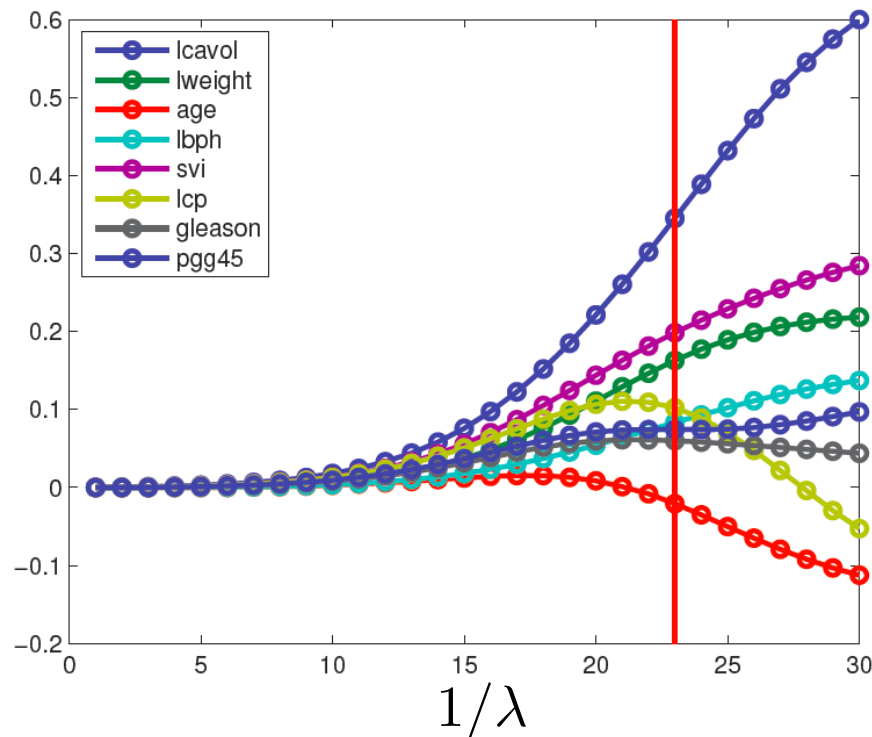
$$\hat{w}_{ridge} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2$$



# Finding best subset: Regularize

Ridge regression makes coefficients small

$$\hat{w}_{ridge} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2$$

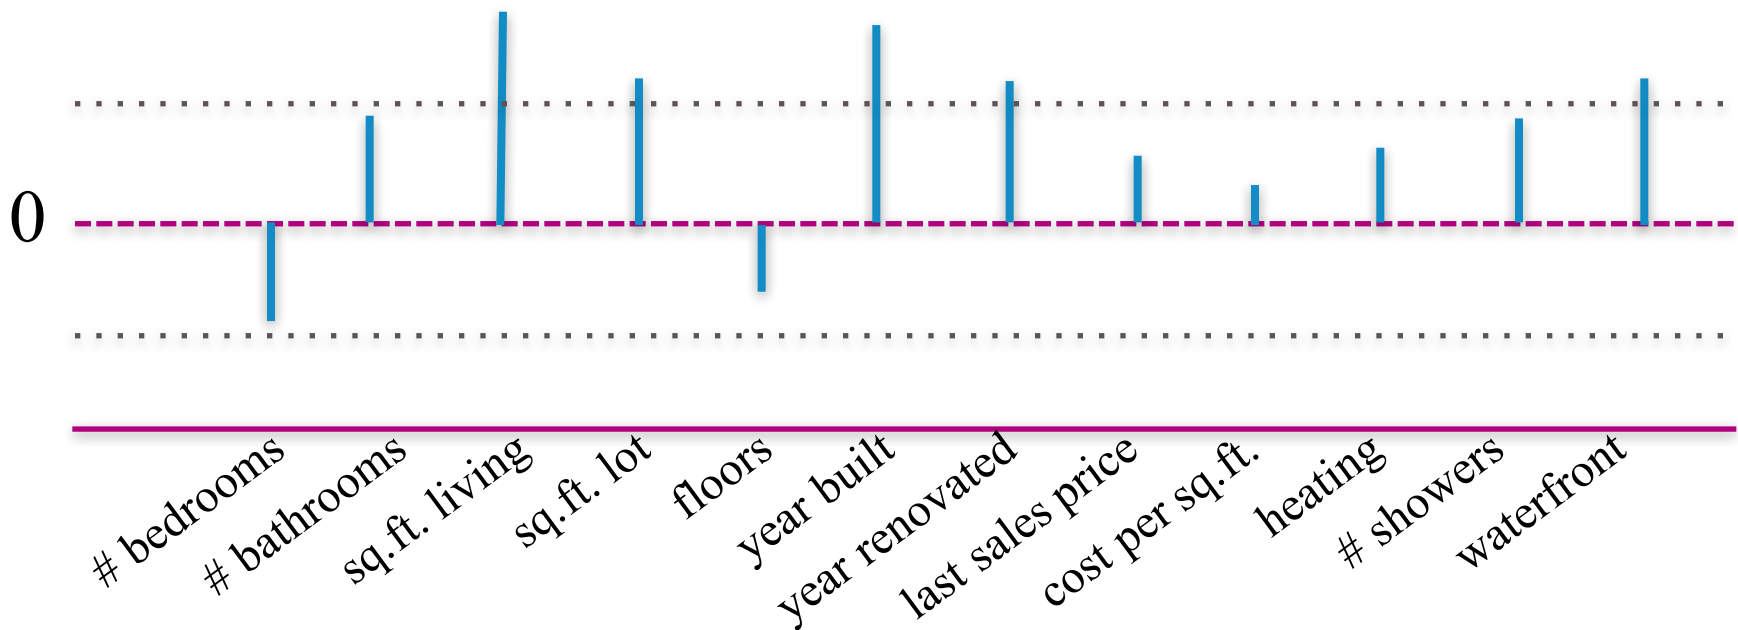


From  
Kevin Murphy  
textbook

# Thresholded Ridge Regression

$$\hat{w}_{ridge} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2$$

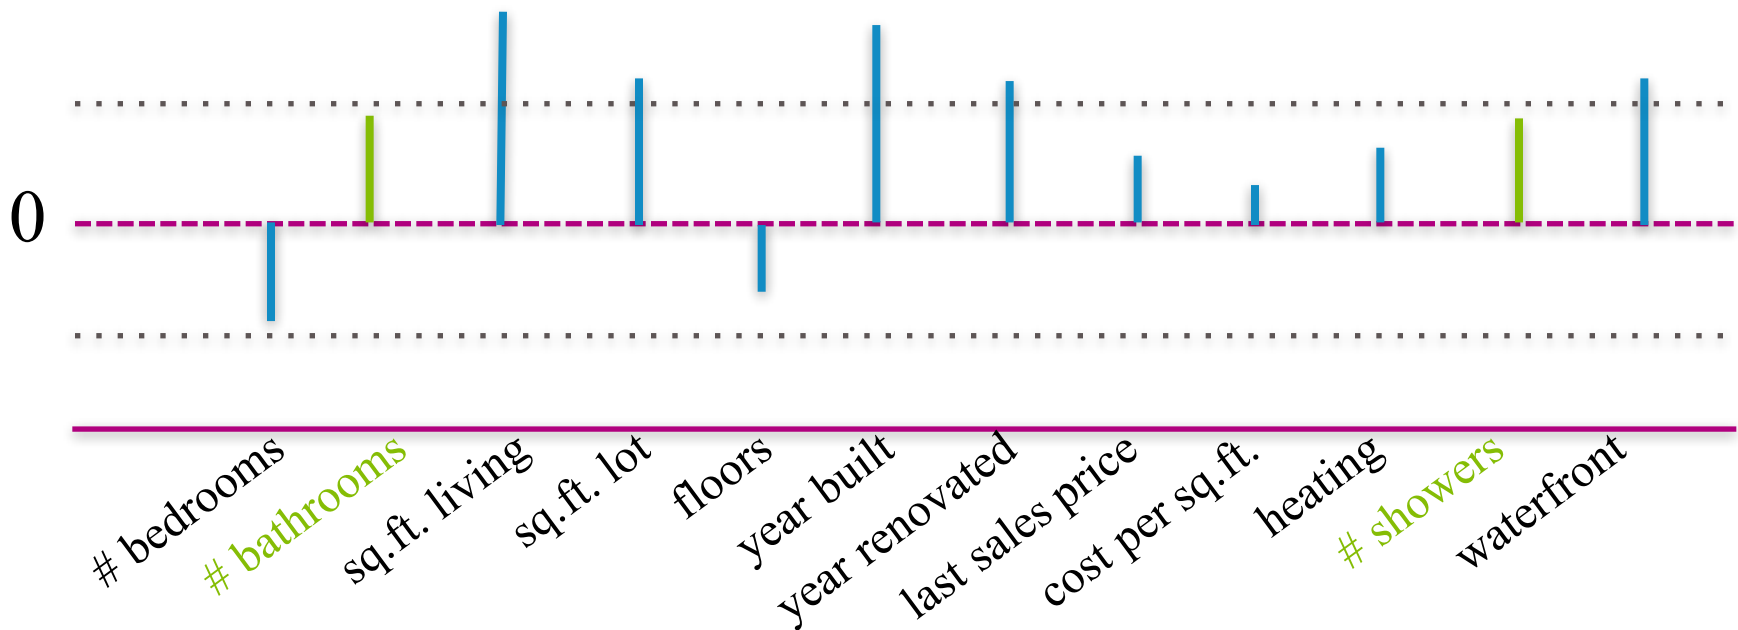
Why don't we just set **small** ridge coefficients to 0?



# Thresholded Ridge Regression

$$\hat{w}_{ridge} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2$$

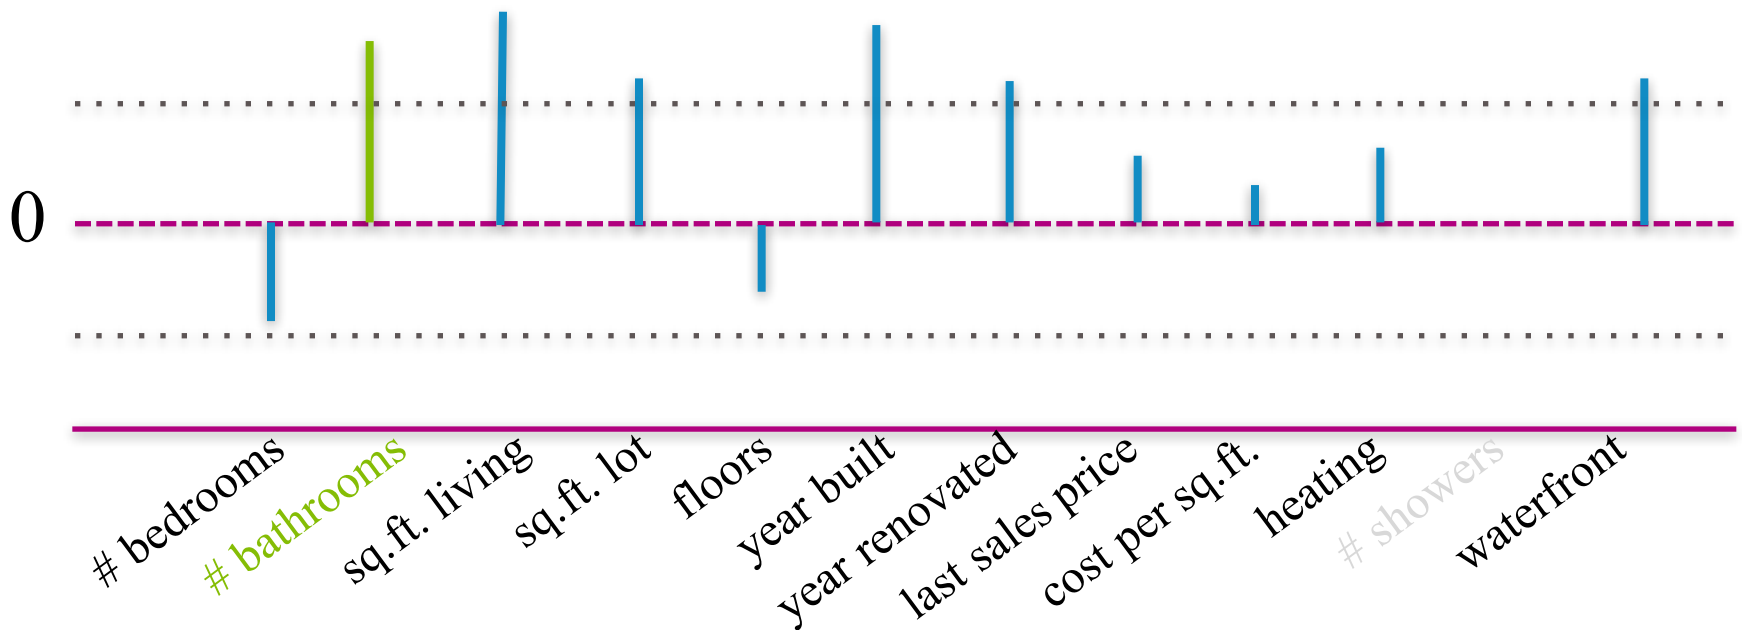
Consider two **related** features (bathrooms, showers)



# Thresholded Ridge Regression

$$\hat{w}_{ridge} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2$$

What if we **didn't** include showers? Weight on bathrooms increases!



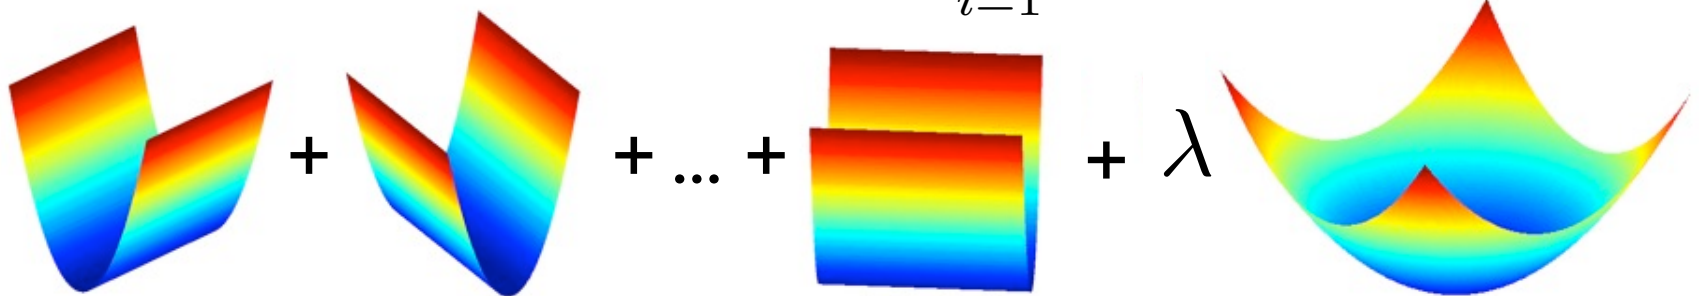
**Can another regularizer perform selection automatically?**



# Recall Ridge Regression

- Ridge Regression objective:

$$\hat{w}_{ridge} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2$$

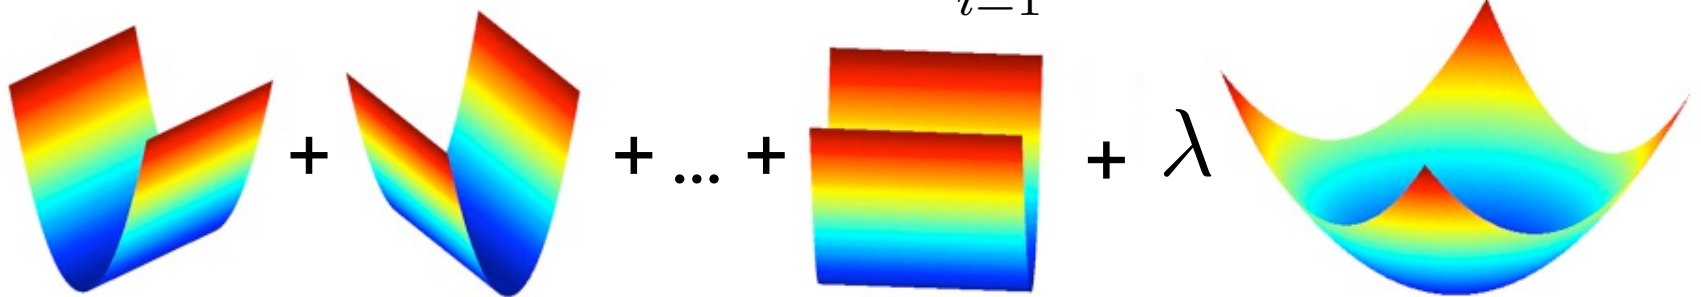


$$||w||_p = \left( \sum_{i=1}^d |w|^p \right)^{1/p}$$

# Ridge vs. Lasso Regression

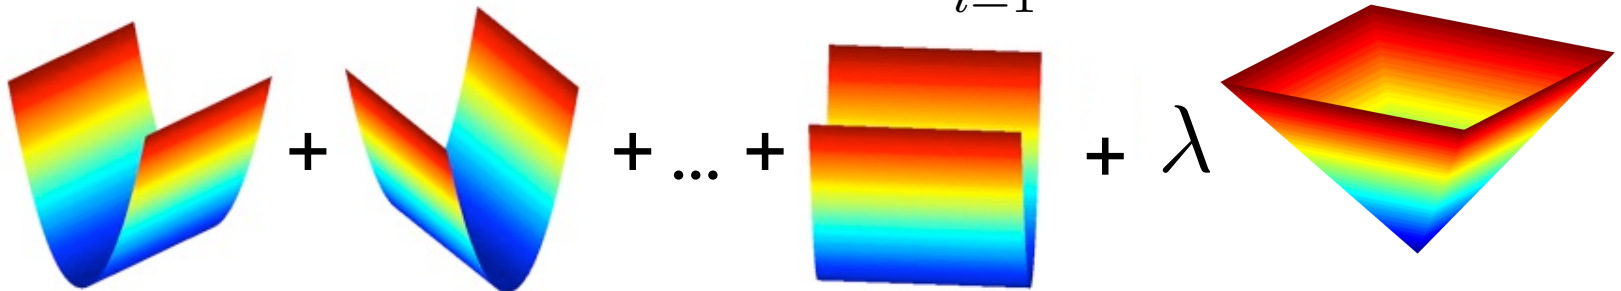
- Ridge Regression objective:

$$\hat{w}_{ridge} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_2^2$$



- Lasso objective:

$$\hat{w}_{lasso} = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda ||w||_1$$



# Penalized Least Squares

Ridge :  $r(w) = ||w||_2^2$       Lasso :  $r(w) = ||w||_1$

$$\hat{w}_r = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda r(w)$$

# Penalized Least Squares

$$\text{Ridge : } r(w) = \|w\|_2^2 \quad \text{Lasso : } r(w) = \|w\|_1$$

$$\hat{w}_r = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda r(w)$$

For any  $\lambda \geq 0$  for which  $\hat{w}_r$  achieves the minimum, there exists a  $\nu \geq 0$  such that

$$\hat{w}_r = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 \quad \text{subject to } r(w) \leq \nu$$

# Penalized Least Squares

Ridge :  $r(w) = ||w||_2^2$

Lasso :  $r(w) = ||w||_1$

$$\hat{w}_r = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 + \lambda r(w)$$

For any  $\lambda \geq 0$  for which  $\hat{w}_r$  achieves the minimum, there exists a  $\nu \geq 0$  such that

$$\hat{w}_r = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 \quad \text{subject to } r(w) \leq \nu$$

